



Figure 3 | Current-split forecasting comparison

Original appearance



Updated data, preserved style




The annotated heatmap style is retained. Rows now contain persistence, control and AGF, the three methods evaluated on the current split; unsupported historical baseline rows are removed.


Left: historical visual reference. Right: figures generated from the current experimental evidence.

Figure 4 | Five-seed primary comparison

Original appearance



Updated data, preserved style




The four panels, blue/orange seed points, black mean diamonds and sample-SD bars are retained. Every point and mean-change label uses the current five-seed data.


Left: historical visual reference. Right: figures generated from the current experimental evidence.

Figure 5 | Forecast errors across horizons

Original appearance



Updated data, preserved style




The four-panel layout, task colours and three curve styles are retained. Curves now summarize the five current AGF seeds with sample-SD bands on all 657 test windows.


Left: historical visual reference. Right: figures generated from the current experimental evidence.

Figure 6 | Comparison under one training protocol

Original appearance





Updated data, preserved style



Grouped bars, control hatching and value labels are retained. The four task panels use valid physical scales, and the comparison no longer claims to isolate a training-schedule effect.

Left: historical visual reference. Right: figures generated from the current experimental evidence.



Figure 7 | Sensitivity to missing input nodes



The grey background, white grid, colours and line patterns are retained. Relative MAE changes use the same 64 test windows, averaging masks within each seed before calculating seed variation.

Left: historical visual reference. Right: figures generated from the current experimental evidence.

Figure 8 | Distribution of learned gate activations



The horizontal distributions, white IQR boxes and mean diamonds are retained. Four tasks use full-test, all-channel gate counts; histogram-based percentiles are explicitly approximate.

Left: historical visual reference. Right: figures generated from the current experimental evidence.