


Forecast sensitivity to missing input nodes



64 fixed test windows · five AGF training seeds · three mask draws per nonzero ratio

Shading: ± 1 sample SD across training seeds after averaging the mask draws